

✖ Importing libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfTransformer #considers how common the word occur in the email
from sklearn.naive_bayes import GaussianNB, MultinomialNB #algorithms for classification tasks
from sklearn.metrics import classification_report, confusion_matrix
```

✖ Reading data

```
df = pd.read_csv('/content/spam_ham_dataset.csv')
df.head()
```

	Unnamed: 0	label	text	label_num
0	605	ham	Subject: enron methanol ; meter # : 988291\r\n...	0
1	2349	ham	Subject: hpl nom for january 9 , 2001\r\n(see...	0
2	3624	ham	Subject: neon retreat\r\nho ho ho , we ' re ar...	0
3	4685	spam	Subject: photoshop , windows , office . cheap ...	1
4	2030	ham	Subject: re : indian springs\r\nthis deal is t...	0

✖ Added name to the first column as "ID"

```
df.rename(columns = {'Unnamed: 0': 'ID'},inplace = True)
```

```
df.head()
```

	ID	label	text	label_num
0	605	ham	Subject: enron methanol ; meter # : 988291\r\n...	0
1	2349	ham	Subject: hpl nom for january 9 , 2001\r\n(see...	0
2	3624	ham	Subject: neon retreat\r\nho ho ho , we ' re ar...	0
3	4685	spam	Subject: photoshop , windows , office . cheap ...	1
4	2030	ham	Subject: re : indian springs\r\nthis deal is t...	0

✖ Generating Profile Report

```
!pip install ydata_profiling
```

```
Collecting ydata_profiling
  Downloading ydata_profiling-4.8.3-py2.py3-none-any.whl (359 kB)
    359.5/359.5 kB 4.2 MB/s eta 0:00:00
Requirement already satisfied: scipy<1.14,>=1.4.1 in /usr/local/lib/python3.10/dist-packages (from ydata_profiling) (1.11.4)
Requirement already satisfied: pandas!=1.4.0,<3,>=1.1 in /usr/local/lib/python3.10/dist-packages (from ydata_profiling) (2.0.3)
Requirement already satisfied: matplotlib<3.9,>=3.2 in /usr/local/lib/python3.10/dist-packages (from ydata_profiling) (3.7.1)
Requirement already satisfied: pydantic>=2 in /usr/local/lib/python3.10/dist-packages (from ydata_profiling) (2.7.1)
Requirement already satisfied: PyYAML<6.1,>=5.0.0 in /usr/local/lib/python3.10/dist-packages (from ydata_profiling) (6.0.1)
Requirement already satisfied: jinja2<3.2,>=2.11.1 in /usr/local/lib/python3.10/dist-packages (from ydata_profiling) (3.1.4)
Collecting visions[type_image_path]<0.7.7,>=0.7.5 (from ydata_profiling)
  Downloading visions-0.7.6-py3-none-any.whl (104 kB)
    104.8/104.8 kB 4.6 MB/s eta 0:00:00
Requirement already satisfied: numpy<2,>=1.16.0 in /usr/local/lib/python3.10/dist-packages (from ydata_profiling) (1.25.2)
Collecting htmlmin==0.1.12 (from ydata_profiling)
```

```

Downloading htmlmin-0.1.12.tar.gz (19 kB)
Preparing metadata (setup.py) ... done
Collecting phik<0.13,>=0.11.1 (from ydata_profiling)
  Downloading phik-0.12.4-cp310-cp310-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (686 kB)
    686.1/686.1 kB 7.2 MB/s eta 0:00:00
Requirement already satisfied: requests<3,>=2.24.0 in /usr/local/lib/python3.10/dist-packages (from ydata_profiling) (2.31.0)
Requirement already satisfied: tqdm<5,>=4.48.2 in /usr/local/lib/python3.10/dist-packages (from ydata_profiling) (4.66.4)
Requirement already satisfied: seaborn<0.14,>=0.10.1 in /usr/local/lib/python3.10/dist-packages (from ydata_profiling) (0.13.1)
Collecting multimethod<2,>=1.4 (from ydata_profiling)
  Downloading multimethod-1.11.2-py3-none-any.whl (10 kB)
Requirement already satisfied: statsmodels<1,>=0.13.2 in /usr/local/lib/python3.10/dist-packages (from ydata_profiling) (0.14.2)
Collecting typeguard<5,>=3 (from ydata_profiling)
  Downloading typeguard-4.2.1-py3-none-any.whl (34 kB)
Collecting imagehash==4.3.1 (from ydata_profiling)
  Downloading ImageHash-4.3.1-py2.py3-none-any.whl (296 kB)
    296.5/296.5 kB 4.8 MB/s eta 0:00:00
Requirement already satisfied: wordcloud>=1.9.1 in /usr/local/lib/python3.10/dist-packages (from ydata_profiling) (1.9.3)
Collecting dacite>=1.8 (from ydata_profiling)
  Downloading dacite-1.8.1-py3-none-any.whl (14 kB)
Requirement already satisfied: numba<1,>=0.56.0 in /usr/local/lib/python3.10/dist-packages (from ydata_profiling) (0.58.1)
Requirement already satisfied: PyWavelets in /usr/local/lib/python3.10/dist-packages (from imagehash==4.3.1->ydata_profiling) (1.6.0)
Requirement already satisfied: pillow in /usr/local/lib/python3.10/dist-packages (from imagehash==4.3.1->ydata_profiling) (9.4.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from jinja2<3.2,>=2.11.1->ydata_profiling) (2.1.2)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib<3.9,>=3.2->ydata_profiling) (1.0.7)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib<3.9,>=3.2->ydata_profiling) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib<3.9,>=3.2->ydata_profiling) (4.22.0)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib<3.9,>=3.2->ydata_profiling) (1.4.5)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib<3.9,>=3.2->ydata_profiling) (23.1)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib<3.9,>=3.2->ydata_profiling) (3.1.0)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib<3.9,>=3.2->ydata_profiling) (2.8.2)
Requirement already satisfied: llvmlite<0.42,>=0.41.0dev0 in /usr/local/lib/python3.10/dist-packages (from numba<1,>=0.56.0->ydata_profiling) (0.42.0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas!=1.4.0,<3,>=1.1->ydata_profiling) (2023.3)
Requirement already satisfied: tzdata>=2022.1 in /usr/local/lib/python3.10/dist-packages (from pandas!=1.4.0,<3,>=1.1->ydata_profiling) (2023.3)
Requirement already satisfied: joblib>=0.14.1 in /usr/local/lib/python3.10/dist-packages (from phik<0.13,>=0.11.1->ydata_profiling) (1.3.2)
Requirement already satisfied: annotated-types>=0.4.0 in /usr/local/lib/python3.10/dist-packages (from pydantic>=2->ydata_profiling) (0.6.0)
Requirement already satisfied: pydantic-core==2.18.2 in /usr/local/lib/python3.10/dist-packages (from pydantic>=2->ydata_profiling) (2.18.2)
Requirement already satisfied: typing-extensions>=4.6.1 in /usr/local/lib/python3.10/dist-packages (from pydantic>=2->ydata_profiling) (4.7.1)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.24.0->ydata_profiling) (3.2.0)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.24.0->ydata_profiling) (3.4)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.24.0->ydata_profiling) (2.0.7)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.24.0->ydata_profiling) (2024.2.2)
Requirement already satisfied: patsy>=0.5.6 in /usr/local/lib/python3.10/dist-packages (from statsmodels<1,>=0.13.2->ydata_profiling) (0.5.6)

```

▼ Profile Report

```

from ydata_profiling import ProfileReport
ProfileReport = ProfileReport(df)
ProfileReport

```

 Summarize dataset: 100%

14/14 [00:03<00:00, 3.62it/s, Completed]

Generate report structure: 100%

1/1 [00:04<00:00, 4.91s/it]

Render HTML: 100%

1/1 [00:00<00:00, 1.36it/s]

Overview

Dataset statistics

Number of variables	4
Number of observations	5171
Missing cells	0
Missing cells (%)	0.0%
Duplicate rows	0
Duplicate rows (%)	0.0%
Total size in memory	161.7 KiB
Average record size in memory	32.0 B

Variable types

Numeric	1
Categorical	2
Text	1


Alerts

ID is highly overall correlated with label1 and 1 other fields (label1, label_num)	High correlation
label1 is highly overall correlated with ID and 1 other fields (ID, label_num)	High correlation

From the data report, we came to know the following things: – Data has no missing values – Each column is highly correlated to with each other. – The number of ham(3672) emails is more than the spam(1499)

▼ Data Statistics

```
df.describe()
```



	ID	label_num
count	5171.000000	5171.000000
mean	2585.000000	0.289886
std	1492.883452	0.453753
min	0.000000	0.000000
25%	1292.500000	0.000000
50%	2585.000000	0.000000
75%	3877.500000	1.000000
max	5170.000000	1.000000

```
df.shape
# 5171 rows and 4 columns
```

```
(5171, 4)
```

✓ Finding Nunique values

```
df.nunique()
```

```
ID          5171
label         2
text        4993
label_num     2
dtype: int64
```


✓ Adding a column length to check the length of messages

```
df['length'] = df['text'].apply(len)
df.head()
```

```
ID  label  text  label_num  length
0   605   ham  Subject: enron methanol ; meter # : 988291\r\n...    0    327
1  2349   ham  Subject: hpl nom for january 9 , 2001\r\n( see...    0     97
2  3624   ham  Subject: neon retreat\r\nho ho ho , we ' re ar...    0   2524
3  4685  spam  Subject: photoshop , windows , office . cheap ...    1    414
4  2030   ham  Subject: re : indian springs\r\nthis deal is t...    0    336
```

✓ Distribution on column length to check which text has highest length

```
plt.figure(figsize = (10,5))
hist_graph = df['length'].plot(kind = 'hist')
hist_graph.set_xlabel("Length")
hist_graph.set_ylabel("Count")
```



The graph elaborates that there are some messages who length is above 4000 or upto 5000

✓ With data statistics we are going to find the longest message in this dataset

```
df.length.describe()
```

```
count    5171.000000
mean     1048.390447
std      1528.514135
min       11.000000
25%      244.000000
50%      540.000000
75%     1237.000000
max     32258.000000
Name: length, dtype: float64
```

The highest message length is 32258 which clearly explains that it is a spam, because spam messages tend to have more text than ham data.

Pre Processing

✓ Removing Punctuations

```
import string
text = df['text']
translator = str.maketrans('', '', string.punctuation)
## using maketrans method to gather punctuations
df['text'] = text.apply(lambda x: x.translate(translator))
#applying translate() to remove punctuations
df.head()
```

	ID	label	text	label_num	length
0	605	ham	Subject enron methanol meter 988291\r\nthis...	0	327
1	2349	ham	Subject hpl nom for january 9 2001\r\n see at...	0	97
2	3624	ham	Subject neon retreat\r\nho ho ho we re aroun...	0	2524
3	4685	spam	Subject photoshop windows office cheap mai...	1	414
4	2030	ham	Subject re indian springs\r\nthis deal is to ...	0	336

✓ Data Test and Train split

```
X_train, X_test, y_train, y_test = train_test_split(df['text'], df['label'], test_size=0.3, random_state=42)
# 30% data for testing and 70% for training
```

✓ Vectorization

```
cv = CountVectorizer() # converts the words into tokenization, converts raw text into numerical representation that machine learning underst
X_train = cv.fit_transform(X_train)
X_test = cv.transform(X_test)
```

✓ TF-IDF

```
tfidf_transformer = TfidfTransformer()
X_train = tfidf_transformer.fit_transform(X_train)
X_test = tfidf_transformer.transform(X_test)
# training and transforming the training dataset into TF-IDF
```

✓ Training model on Naive Bayes

```
Gnb = GaussianNB() #ensures that the probability calculated is normal
Mnb = MultinomialNB()
Gnb_model = Gnb.fit(X_train.toarray(), y_train) #The toarray() method is used on sparse matrices((when there are many 0s in model) in scikit
Mnb_model = Mnb.fit(X_train.toarray(), y_train)
```

✓ Prediction

```
print("Prediction for Gaussian Naive Bayes:", Gnb_model.predict(X_test.toarray()))
```

```
➤ Predicted: ['ham' 'spam' 'ham' ... 'ham' 'spam' 'ham']
```

```
print("Prediction for Multinomial Naive Bayes:", Mnb_model.predict(X_test.toarray()))
```

```
➤ Prediction for Multinomial Naive Bayes: ['ham' 'spam' 'ham' ... 'ham' 'ham' 'ham']
```

Classification Report

✓ Gaussian Naivr Bayes

```
print(classification_report(y_test, Gnb_model.predict(X_test.toarray())))
```

```
➤
```

	precision	recall	f1-score	support
ham	0.96	0.97	0.97	1121
spam	0.92	0.90	0.91	431
accuracy			0.95	1552
macro avg	0.94	0.94	0.94	1552
weighted avg	0.95	0.95	0.95	1552

✓ Multinomial Algorithm

```
print(classification_report(y_test, Mnb_model.predict(X_test.toarray())))
```

```
➤
```

	precision	recall	f1-score	support
ham	0.85	1.00	0.92	1121
spam	1.00	0.52	0.69	431
accuracy			0.87	1552
macro avg	0.92	0.76	0.80	1552
weighted avg	0.89	0.87	0.85	1552

Overall the accuracy of Gaussian Naive Bayes is more than Multinomial Bayes i.e 95% but we cannot rely on accuracy only so moving forward to analyze th confusion matrix


✓ Confusion Matrix for Gnb

```
confusion_matrix(y_test, Gnb_model.predict(X_test.toarray()))
```

```
➤ array([[1088, 33],
        [ 41, 390]])
```

✓ Plotting Confusion matrix

```
plt.figure(figsize = ( 8,4))
sns.heatmap(confusion_matrix(y_test, Gnb_model.predict(X_test.toarray()))), annot = True, fmt = 'd', cmap = 'Greens')
plt.xlabel('Predicted Labels')
plt.ylabel('Actual Labels')
plt.title('Confusion Matrix for Gaussian Naive Bayes')
plt.show()
```



- 1088 instances are correctly classified as "Ham"(True Negative)
- 33 instances were incorrectly classified as "Ham (False Positive)
- 41 were incorrectly classified "Spam"
- 390 were correctly classified "Spam"

✓ Confusion matrix for multinomial bayes

```
confusion_matrix(y_test, Mnb_model.predict(X_test.toarray()))
```

```
array([[1121,  0],
       [ 205, 226]])
```

✓ Plotting

```
plt.figure(figsize= (8,4))
sns.heatmap(confusion_matrix(y_test, Mnb_model.predict(X_test.toarray()))), annot = True, fmt = 'd', cmap = 'Blues')
plt.xlabel('Predicted Labels')
plt.ylabel('Actual Labels')
plt.title('Confusion Matrix for Multinomial Naive Bayes')
plt.show()
```

